

1BM22CS157_BDA_lab7

```
WCDriver.java x WCMapper.java WReducer.java
1 *import java.io.IOException;
12 public class WCDriver extends Configured implements Tool {
13 public int run(String args[]) throws IOException
14 {
15     if (args.length < 2)
16     {
17         System.out.println("Please give valid inputs");
18         return -1;
19     }
20     JobConf conf = new JobConf(WCDriver.class);
21     FileInputFormat.setInputPaths(conf, new Path(args[0]));
22     FileOutputFormat.setOutputPath(conf, new Path(args[1]));
23     conf.setMapperClass(WCMapper.class);
24     conf.setReducerClass(WReducer.class);
25     conf.setMapOutputKeyClass(Text.class);
26     conf.setMapOutputValueClass(IntWritable.class);
27     conf.setOutputKeyClass(Text.class); conf.setOutputValueClass(IntWritable.class);
28     JobClient.runJob(conf);
29     return 0;
30 }
31 // Main Method
32 public static void main(String args[]) throws Exception
33 {
34     int exitCode = ToolRunner.run(new WCDriver(), args);
35     System.out.println(exitCode);
36 }
37 }
```

Console x Problems Debug Shell Coverage

```
WCDriver.java  WCMapper.java x  WCReducer.java
1 import java.io.IOException;
9 public class WCMapper extends MapReduceBase implements Mapper<LongWritable,Text,Text,IntWritable> {
10 // Map function
11 public void map(LongWritable key, Text value, OutputCollector<Text,IntWritable> output, Reporter rep) throws IOException
12 {
13     String line = value.toString();
14     // Splitting the line on spaces
15     for (String word : line.split(" "))
16     {
17         if (word.length() > 0)
18         {
19             output.collect(new Text(word), new IntWritable(1));
20         }
21     }
22 }
23 }
```

Console x Problems Debug Shell Coverage

```
WCDriver.java  WCMapper.java  WCReducer.java ×  
1 *import java.io.IOException;  
9 public class WCReducer extends MapReduceBase implements Reducer<Text,  
10 IntWritable, Text, IntWritable> {  
11 // Reduce function  
12 public void reduce(Text key, Iterator<IntWritable> value, OutputCollector<Text, IntWritable> output, Reporter rep) throws IOException  
13 {  
14     int count = 0;  
15     // Counting the frequency of each words  
16     while (value.hasNext())  
17     {  
18         IntWritable i = value.next();  
19         count += i.get();  
20     }  
21     output.collect(key, new IntWritable(count));  
22 }  
23 }  
24  
  
Console ×  Problems  Debug Shell  Coverage
```

```

hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hadoop jar WordCount.jar WCDriver /shurath/test.txt /shurath/out.txt
2025-04-29 15:11:42,145 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-04-29 15:11:42,185 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-04-29 15:11:42,185 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-04-29 15:11:42,191 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:11:42,254 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:11:42,303 INFO mapred.FileInputFormat: Total input files to process : 1
2025-04-29 15:11:42,328 INFO mapreduce.JobSubmitter: number of splits:1
2025-04-29 15:11:42,390 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1048437123_0001
2025-04-29 15:11:42,390 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:11:42,448 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-04-29 15:11:42,449 INFO mapreduce.Job: Running job: job_local1048437123_0001
2025-04-29 15:11:42,449 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-04-29 15:11:42,451 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2025-04-29 15:11:42,453 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:11:42,453 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:11:42,507 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-04-29 15:11:42,508 INFO mapred.LocalJobRunner: Starting task: attempt_local1048437123_0001_m_000000_0
2025-04-29 15:11:42,519 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:11:42,519 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:11:42,526 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-04-29 15:11:42,529 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/shurath/test.txt:0+89
2025-04-29 15:11:42,538 INFO mapred.MapTask: numReduceTasks: 1
2025-04-29 15:11:42,569 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2025-04-29 15:11:42,569 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-04-29 15:11:42,569 INFO mapred.MapTask: soft limit at 83886080
2025-04-29 15:11:42,569 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-04-29 15:11:42,569 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-04-29 15:11:42,571 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-04-29 15:11:42,608 INFO mapred.LocalJobRunner:
2025-04-29 15:11:42,608 INFO mapred.MapTask: Starting flush of map output
2025-04-29 15:11:42,608 INFO mapred.MapTask: Spilling map output
2025-04-29 15:11:42,608 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufvoid = 104857600
2025-04-29 15:11:42,608 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length = 77/6553600
2025-04-29 15:11:42,611 INFO mapred.MapTask: Finished spill 0
2025-04-29 15:11:42,616 INFO mapred.Task: Task:attempt_local1048437123_0001_m_000000_0 is done. And is in the process of committing
2025-04-29 15:11:42,618 INFO mapred.LocalJobRunner: hdfs://localhost:9000/shurath/test.txt:0+89
2025-04-29 15:11:42,618 INFO mapred.Task: Task 'attempt_local1048437123_0001_m_000000_0' done.
2025-04-29 15:11:42,621 INFO mapred.Task: Final Counters for attempt_local1048437123_0001_m_000000_0: Counters: 23
File System Counters
  FILE: Number of bytes read=4049
  FILE: Number of bytes written=645512
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0

```

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hadoop fs -ls /shuraih/out.txt
Found 2 items
-rw-r--r-- 1 hadoop supergroup      0 2025-04-29 15:11 /shuraih/out.txt/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    69 2025-04-29 15:11 /shuraih/out.txt/part-00000
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$ hadoop fs -cat /shuraih/out.txt/part-00000
are      1
brother  1
family   1
hi       1
how      5
is       4
job      1
sister   1
you      1
your     4
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~/Desktop$
```

```
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmscecse-HP-Elite-Tower-800-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ su hduser
su: user hduser does not exist or the user entry does not contain all the required fields
hadoop@bmscecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ jps
7890 ResourceManager
8054 NodeManager
8488 Jps
7259 DataNode
5468 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734.jar
7117 NameNode
7517 SecondaryNameNode
```