

Hadoop Average and Mean temperature

1BM22CS157

```

bytes written=8
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -ls /shuraih
Found 4 items
drwxr-xr-x - hadoop supergroup      0 2025-04-29 15:11 /shuraih/out.txt
drwxr-xr-x - hadoop supergroup      0 2025-05-06 14:50 /shuraih/output
-rw-r--r-- 1 hadoop supergroup 888190 2025-05-06 14:14 /shuraih/temperature.txt
-rw-r--r-- 1 hadoop supergroup   89 2025-04-15 14:43 /shuraih/test.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -ls /shuraih/output
Found 2 items
-rw-r--r-- 1 hadoop supergroup      0 2025-05-06 14:50 /shuraih/output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    8 2025-05-06 14:50 /shuraih/output/part-r-000000
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -cat /shuraih/output/part-r-000000
1901      46
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -ls /home/hadoop/Desktop/MapMax-ias-MapMaxDriver-/shuraih/temperature.txt /shuraih/output1
bytes written=74
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -ls /shuraih/output1
Found 2 items
-rw-r--r-- 1 hadoop supergroup      0 2025-05-06 15:08 /shuraih/output1/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    74 2025-05-06 15:08 /shuraih/output1/part-r-000000
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop fs -cat /shuraih/output1/part-r-000000
01      4
02      0
03      7
04     44
05     100
06     168
07     219
08     198
09     141
10     100
11      19
12       3
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC: $

```

```

hadoop@bmscece-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop jar /home/hadoop/Desktop/MeanMaxTemp.jar MeanMaxDriver /shurath/temperature.txt /shurath/output1
2025-05-06 15:08:42,166 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 15:08:42,207 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 15:08:42,207 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 15:08:42,268 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-05-06 15:08:42,320 INFO Input.FileInputFormat: Total input files to process : 1
2025-05-06 15:08:42,349 INFO mapreduce.JobSubmitter: number of splits:1
2025-05-06 15:08:42,414 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local2013736307_0001
2025-05-06 15:08:42,414 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 15:08:42,485 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 15:08:42,485 INFO mapreduce.Job: Running job: job_local2013736307_0001
2025-05-06 15:08:42,486 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-05-06 15:08:42,489 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:08:42,489 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:08:42,489 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:08:42,490 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 15:08:42,539 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-05-06 15:08:42,539 INFO mapred.LocalJobRunner: Starting task: attempt_local2013736307_0001_m_000000_0
2025-05-06 15:08:42,550 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 15:08:42,550 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 15:08:42,550 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 15:08:42,557 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-05-06 15:08:42,559 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/shurath/temperature.txt:0+888190
2025-05-06 15:08:42,595 INFO mapred.MapTask: (EQUATOR) 0 kv: 26214396(104857504)
2025-05-06 15:08:42,595 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 15:08:42,595 INFO mapred.MapTask: soft limit at 83886080
2025-05-06 15:08:42,595 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-05-06 15:08:42,595 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-05-06 15:08:42,597 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-05-06 15:08:42,667 INFO mapred.LocalJobRunner:
2025-05-06 15:08:42,668 INFO mapred.MapTask: Starting flush of map output
2025-05-06 15:08:42,668 INFO mapred.MapTask: Spilling map output
2025-05-06 15:08:42,668 INFO mapred.MapTask: bufstart = 0; bufend = 45948; bufvoid = 104857600
2025-05-06 15:08:42,668 INFO mapred.MapTask: kvstart = 26214396(104857504); kvend = 26188144(104752576); length = 26253/6553600
2025-05-06 15:08:42,678 INFO mapred.MapTask: Finished spill 0
2025-05-06 15:08:42,684 INFO mapred.Task: Task:attempt_local2013736307_0001_m_000000_0 is done. And is in the process of committing
2025-05-06 15:08:42,685 INFO mapred.LocalJobRunner: map
2025-05-06 15:08:42,686 INFO mapred.Task: Task 'attempt_local2013736307_0001_m_000000_0' done.
2025-05-06 15:08:42,688 INFO mapred.Task: Final Counters for attempt_local2013736307_0001_m_000000_0: Counters: 23
File System Counters
  FILE: Number of bytes read=4367
  FILE: Number of bytes written=703883
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=888190
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=5
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=1
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=6565
  Map output records=6564
  Map output bytes=45948

```

```
hadoop@bmsccsc-HP-Elite-Tower-800-G9-Desktop-PC: $ hadoop jar /home/hadoop/Desktop/TemperatureAvg.jar AverageDriver /shurath/temperature.txt /shurath/output
2025-05-06 14:50:01,243 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-05-06 14:50:01,284 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-05-06 14:50:01,284 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-05-06 14:50:01,352 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-05-06 14:50:01,407 INFO Input.FileInputFormat: Total input files to process : 1
2025-05-06 14:50:01,468 INFO mapreduce.JobSubmitter: number of splits:1
2025-05-06 14:50:01,542 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local251162856_0001
2025-05-06 14:50:01,542 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-05-06 14:50:01,608 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-05-06 14:50:01,609 INFO mapreduce.Job: Running job: job_local251162856_0001
2025-05-06 14:50:01,609 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-05-06 14:50:01,612 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 14:50:01,613 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 14:50:01,613 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 14:50:01,614 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2025-05-06 14:50:01,780 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-05-06 14:50:01,781 INFO mapred.LocalJobRunner: Starting task: attempt_local251162856_0001_m_000000_0
2025-05-06 14:50:01,804 INFO output.PathOutputCommitterFactory: No output committer factory defined, defaulting to FileOutputCommitterFactory
2025-05-06 14:50:01,804 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-05-06 14:50:01,804 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2025-05-06 14:50:01,811 INFO mapred.Task: Using ResourceCalculatorProcessTree: [ ]
2025-05-06 14:50:01,813 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/shurath/temperature.txt:0+888190
2025-05-06 14:50:01,849 INFO mapred.MapTask: (EQUATOR) 0 kvl 26214396(104857584)
2025-05-06 14:50:01,849 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2025-05-06 14:50:01,849 INFO mapred.MapTask: soft limit at 83886080
2025-05-06 14:50:01,849 INFO mapred.MapTask: bufstart = 0; bufvold = 104857600
2025-05-06 14:50:01,849 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-05-06 14:50:01,851 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-05-06 14:50:01,991 INFO mapred.LocalJobRunner:
2025-05-06 14:50:01,991 INFO mapred.MapTask: Starting flush of map output
2025-05-06 14:50:01,992 INFO mapred.MapTask: Spilling map output
2025-05-06 14:50:01,992 INFO mapred.MapTask: bufstart = 0; bufend = 59076; bufvold = 104857600
2025-05-06 14:50:01,992 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26188144(104752576); length = 26253/6553600
2025-05-06 14:50:02,000 INFO mapred.MapTask: Finished spill 0
2025-05-06 14:50:02,005 INFO mapred.Task: Task:attempt_local251162856_0001_m_000000_0 is done. And is in the process of committing
2025-05-06 14:50:02,007 INFO mapred.LocalJobRunner: map
2025-05-06 14:50:02,007 INFO mapred.Task: Task 'attempt_local251162856_0001_m_000000_0' done.
2025-05-06 14:50:02,009 INFO mapred.Task: Final Counters for attempt_local251162856_0001_m_000000_0: Counters: 23
  File System Counters
    FILE: Number of bytes read=4258
    FILE: Number of bytes written=713814
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=888190
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=5
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=1
    HDFS: Number of bytes read erasure-coded=0
  Map-Reduce Framework
    Map input records=6564
    Map output records=6564
    Map output bytes=59076
```